
DETERMINISTIC MEMORY FUSION FOR LONG-HORIZON CONVERSATIONAL QA: CROSS-MODEL EVIDENCE BEYOND LLM-GUIDED FUSION

ABSTRACT

To maintain factual coherence across long conversational sessions, memory-augmented agents must continuously consolidate redundant dialogue fragments into compact, retrievable representations — a process known as memory fusion that accounts for the dominant share of system performance, with its removal collapsing multi-hop F1 by over 53%. However, fusion universally relies on additional LLM calls at maintenance time, a dependency whose necessity has never been empirically tested. We ask a precise question: is the fusion gain driven by the LLM’s linguistic rewriting, or by the simpler effect of deduplicating and packing verbatim evidence into a fixed retrieval budget? To answer it, we propose **DFM-Fusion**, a fully deterministic operator requiring zero maintenance-time LLM calls. More concretely, DFM-Fusion applies four sequential stages — sentence segmentation, near-duplicate removal, MMR-based sentence packing, and deterministic coverage verification — to consolidate memory clusters without any generative rewriting. Key to DFM-Fusion’s design is the recognition that verbatim evidence packing, not LLM rewriting, accounts for virtually all of the fusion gain in fact-centric conversational QA. Across six open-weight model families on LoCoMo, DFM-Fusion achieves the best average multi-hop F1 (0.1877) versus LLM-Fusion (0.1841) and No-Fusion (0.1785), improving in 5/6 models while reducing maintenance overhead by $69.13\times$ and eliminating all fusion-specific API calls. We further propose **DFM-Route**, a query-type-aware retrieval extension requiring no additional LLM calls, which addresses the temporal gap that verbatim packing alone cannot close: across six models, temporal F1 improves $5.5\times$ (0.055 to 0.298), and DFM-Route with DeepSeek-V3.2 reaches 0.4969 overall F1 — the highest result across all experiments. Together, DFM-Fusion and DFM-Route demonstrate that the full memory pipeline, from consolidation to evidence delivery, can operate without any API dependency.

1 INTRODUCTION

Consider a long-running personal assistant that has conversed with a user across hundreds of sessions over months. When the user asks a multi-hop question — “Which hobby did I pick up right after quitting the gym?” — the agent must integrate two distant conversation fragments out of thousands of accumulated memory entries. The challenge is not retrieval quality alone: as related interactions accumulate, a fixed retrieval budget increasingly fills with near-duplicate evidence rather than the distinct facts that such questions demand. This retrieval-noise problem is not hypothetical: on LoCoMo’s multi-hop benchmark, which requires integrating evidence across multiple conversation sessions, even strong systems score below 20 F1 without effective memory consolidation (Maharana et al., 2024).

Memory fusion — the consolidation of semantically similar fragments into compact, non-redundant representations — is the single largest performance driver in state-of-the-art long-horizon memory systems. In FadeMem’s ablation study, removing fusion while holding everything else fixed drops multi-hop F1 from 29.43 to 13.63, a 53.7% collapse (Wei et al., 2026). Fusion is not a peripheral optimization; it is what makes long-horizon conversational QA work at all.

However, every system that performs fusion does so through LLM-guided consolidation: an LLM rewrites the memory cluster into a coherent summary, then a second LLM call verifies that the rewrite preserved information (Wei et al., 2026; Chhikara et al., 2025). This pipeline works, but it rests on an unexamined assumption whose necessity has never been empirically tested. Our key insight is that in fact-centric conversational QA, the fusion gain comes almost entirely from deduplicating and packing verbatim evidence into the fixed retrieval budget — not from the LLM’s linguistic rewriting ability. Different from prior works that treat LLM rewriting as an indispensable step, we show that a fully deterministic operator preserving verbatim sentences recovers the same gain while requiring zero additional LLM calls at maintenance time. This changes how practitioners design agent memory: persistent, consolidating memory becomes deployable without any API dependency in the maintenance loop, making it accessible for local deployments, cost-sensitive applications, and settings where reproducibility or auditability is required.

We propose **DFM-Fusion** (Deterministic FadeMem-Fusion), a four-stage pipeline that replaces LLM-guided fusion with fully deterministic operations. Concretely, given a cluster of semantically similar memory entries, DFM-Fusion applies sentence segmentation, near-duplicate removal via embedding similarity, MMR-based sentence packing within a token budget, and salient-token coverage verification with a fallback guarantee — without invoking any LLM at any stage. We evaluate the same memory pipeline across six open-weight model families on LoCoMo while holding the retrieval budget, scoring protocol, and answer-generation stack fixed. Our contributions are:

- A deterministic, quote-preserving fusion operator that eliminates all LLM calls from memory consolidation while maintaining auditable, verbatim evidence packing.
- A six-model family evaluation on LoCoMo showing that DFM-Fusion achieves the strongest average multi-hop F1 (0.1877) and improves over LLM-Fusion on the primary metric in 5/6 models — confirming that the fusion gain is driven by information packing, not LLM rewriting.
- Consistent operational gains: $69.13\times$ average maintenance speedup, 11.74% reduction in total API calls, and complete elimination of fusion-specific LLM calls — with no performance cost.

The conclusion is deliberately scoped: for fact-centric long-horizon conversational QA, where the evaluation rewards verbatim evidence retrieval, LLM-guided fusion is unnecessary. Whether the same holds for tasks requiring abstraction or conflict reconciliation remains open, and we discuss this boundary explicitly in Section 5.

2 RELATED WORK

Memory Systems for Conversational Agents. Long-horizon conversational agents require persistent memory to maintain coherence across extended interactions. Earlier work established the core architectural patterns: MemGPT (Packer et al., 2023) pioneered hierarchical storage tiers through an operating-system metaphor for LLM context management, while FadeMem (Wei et al., 2026) introduced biologically-inspired forgetting with LLM-guided fusion that consolidates related memories as less relevant ones decay. More recent systems extend this foundation through alternative organizational schemes: Mem0 (Chhikara et al., 2025) captures relational structure via graph-based representations, TiMem (Li et al., 2026) organizes memories through temporal hierarchies, SimpleMem (Liu et al., 2026) achieves efficiency through semantic compression, and EverMemOS (Hu et al., 2026) models engram-inspired memory lifecycles. However, across all of these systems, memory consolidation invariably requires additional LLM calls — for rewriting, merging, or preservation verification — a dependency whose necessity has never been empirically isolated from the simpler effect of redundancy control under a fixed retrieval budget.

Redundancy-Aware Selection in Information Retrieval. The algorithmic foundations for deterministic redundancy control are well established in information retrieval. Maximal Marginal Relevance (Carbonell & Goldstein, 1998) explicitly trades off relevance and diversity during document ranking, and near-duplicate removal via embedding similarity is a standard preprocessing step in retrieval pipelines. However, these techniques operate at query time — selecting among already-stored documents to populate a retrieval result — rather than at maintenance time, where the task

is to consolidate redundant memory entries into compact representations before retrieval occurs. Whether their principles transfer directly to memory maintenance, replacing generative rewriting, has not been tested.

Retrieval-Augmented Generation. RAG approaches (Gao et al., 2023) retrieve relevant context to augment LLM generation. HippoRAG (Gutierrez et al., 2024) draws inspiration from hippocampal indexing theory, using knowledge graphs and personalized PageRank for retrieval. GraphRAG (Edge et al., 2024) leverages entity knowledge graphs with community summaries for query-focused summarization. While these methods advance retrieval quality, they assume a static document store and do not address the memory consolidation problem where redundant entries must be continuously merged while preserving information fidelity.

Memory Benchmarks. LoCoMo (Maharana et al., 2024) evaluates long-term conversational memory through multi-hop question answering over 300-turn dialogues spanning 35 sessions. BEAM (Tavakoli et al., 2025) extends evaluation to million-token contexts with diverse memory probing tasks. These benchmarks reveal that even long-context LLMs struggle with extended dialogues, motivating the need for effective memory management — yet neither directly evaluates whether LLM-guided memory consolidation is necessary, or whether equivalent performance can be achieved through deterministic maintenance operators.

3 METHOD

We propose DFM-Fusion (Deterministic FadeMem-Fusion), a fully deterministic replacement for LLM-guided memory fusion in long-horizon conversational agents. Our approach maintains the memory architecture and retrieval mechanisms of existing systems while eliminating all LLM calls from the fusion operator.

3.1 PROBLEM FORMULATION

Memory fusion addresses the challenge of consolidating redundant information in conversational memory stores. Given a cluster of semantically similar memory entries $\mathcal{M} = \{m_1, m_2, \dots, m_k\}$ identified through temporal-semantic clustering, the fusion task produces a single consolidated entry m_f that preserves essential information within a token budget B .

Existing approaches (Wei et al., 2026; Chhikara et al., 2025) employ LLMs to rewrite and merge memories, then use an additional LLM call to verify that the rewrite preserved information. As discussed in Section 1, this introduces deployment challenges related to cost, determinism, and auditability. Our approach addresses these by preserving verbatim spans while enforcing explicit token budgets.

Memory strength. We follow the FadeMem-style strength dynamics used by the implementation. Each new memory starts with strength $v_i = 1.0$ in the short-term layer. At time t , its importance is

$$I_i(t) = \alpha r_i(t) + \beta \frac{a_i}{1 + a_i} + \gamma \exp(-\delta \Delta t_i), \quad (1)$$

where $r_i(t)$ is cosine relevance to the current session context when available, a_i is access frequency, and Δt_i is age in days. Strength decays as

$$v_i(t) = v_i(t_0) \exp[-\lambda_{\text{base}} \exp(-\mu I_i(t)) (\Delta t_i)^\rho], \quad (2)$$

with $\rho = 0.8$ for long-term memories and $\rho = 1.2$ for short-term memories. Access consolidation increases strength by

$$v_i \leftarrow v_i + d_v(1 - v_i) \exp(-c_i/N), \quad (3)$$

where c_i is access count. A fused memory receives strength

$$v_f = \min \left(1, \max_i v_i + \epsilon \text{Var}(\{v_i\}) \right), \quad (4)$$

and is inserted into the long-term layer. In our experiments $\alpha = 0.4$, $\beta = 0.3$, $\gamma = 0.3$, $\delta = 0.1$, $\mu = 2.0$, $d_v = 0.15$, $N = 5$, and $\epsilon = 0.1$.

3.2 DFM-FUSION PIPELINE

Given a cluster $\mathcal{M}$ exceeding a minimum size threshold, DFM-Fusion produces fused text s_f through four deterministic stages. More specifically: (1) sentence segmentation breaks each memory entry into sentence-level units to enable fine-grained selection; (2) near-duplicate removal eliminates redundant sentences via embedding similarity; (3) MMR-based sentence packing selects a diverse, high-strength subset within the token budget; and (4) coverage verification confirms that salient evidence is preserved, with a fallback guarantee if it is not. No LLM is invoked at any stage.

Stage 1: Sentence Segmentation. Each memory entry m_i is split into individual sentences using rule-based tokenization. However, applying fusion directly at the memory-entry level would truncate content mid-sentence when the token budget is enforced, fragmenting the verbatim spans that downstream retrieval and F1 evaluation depend on. Sentence-level segmentation ensures that every selected unit is a complete, citable span.

Stage 2: Near-Duplicate Removal. We compute pairwise cosine similarity between sentence embeddings using Sentence-BERT (Reimers & Gurevych, 2019) (384-dim). However, the raw sentence pool from a cluster frequently contains near-identical reformulations of the same fact — variants that would consume token budget without adding evidence coverage. Sentences with cosine similarity exceeding threshold $\theta_{\text{dup}} = 0.85$ are treated as near-duplicates; we retain the sentence from the memory with higher strength $v_i(t)$, using a more recent timestamp as a tie-breaker. The same Sentence-BERT encoder is reused throughout the pipeline — for deduplication, for MMR scoring, and for computing the retrieval embedding of the final fused entry — eliminating any inconsistency between maintenance-time and retrieval-time representations.

Stage 3: MMR-Based Sentence Packing. From the deduplicated sentence pool, we select sentences using a strength-aware variant of Maximal Marginal Relevance (Carbonell & Goldstein, 1998) with diversity parameter $\lambda = 0.7$. However, relevance scoring alone would fill the budget with the highest-strength sentences from the dominant source memory, crowding out lower-strength entries that may contain complementary evidence. The MMR objective balances source-memory strength against redundancy with already-selected sentences:

$$\text{MMR}(s_i) = \lambda \cdot \hat{v}(s_i) - (1 - \lambda) \cdot \max_{s_j \in S} \cos(s_i, s_j) \quad (5)$$

where $\hat{v}(s_i) \in [0, 1]$ is the normalized strength of the source memory for sentence s_i , and S is the set of already-selected sentences. We greedily pick $\arg \max_{s_i \in R \setminus S} \text{MMR}(s_i)$ until adding the next sentence would exceed the token budget $B_{\text{fuse}} = 768$ tokens. Selected sentences are sorted by source-memory timestamp before concatenation. The fused memory embedding is computed from the actual fused text s_f , not by averaging constituent embeddings; this ensures retrieval operates on a representation that accurately reflects the consolidated content, and was validated by our ablation (Table 1, lower block).

Stage 4: Coverage Verification. To confirm information preservation without an LLM, we implement a deterministic coverage check. We extract salient evidence from the original cluster: numbers and date-like expressions, capitalized entities, top- K TF-IDF weighted tokens ($K = 20$), and stemmed n -gram coverage features in the enriched configuration. Coverage recall is the fraction of these tokens appearing in s_f . It should be noted that a coverage gate that simply rejects low-recall fusions would leave clusters unmerged, breaking the consolidation guarantee. To address this, when recall falls below threshold $\theta_{\text{cov}} = 0.85$, we fall back to concatenating the highest-strength original memories within the budget in strength-descending order, so a fused result is always emitted rather than leaving the cluster unmerged.

3.3 DFM-ROUTE: QUERY-AWARE RETRIEVAL EXTENSION

While DFM-Fusion operates at maintenance time — consolidating memory clusters offline before any query arrives — DFM-Route extends the system at retrieval time by routing each query to a retrieval strategy matched to its evidence requirements. Like DFM-Fusion, DFM-Route introduces zero LLM calls at any stage; all routing and verification decisions are fully deterministic.

Query Type Classification. Each incoming query is classified into one of eight types by a rule-based regex classifier: *temporal* (when-questions), *list*, *number*, *who*, *where*, *yes_no*, *inferential_yes_no*, or *fact* (default). Classification is based solely on surface patterns (e.g., leading “when”, “who”, “how many”) without any learned model.

Evidence Gate Score. For each query, two candidate retrieval results are generated: a *conservative* set (standard cosine-similarity ranking) and an *expansive* set (anchor-then-MMR, which locks on the highest-similarity memory and applies MMR diversification). A deterministic gate score $g \in [0, 1]$ is computed for each candidate set:

$$g = 0.35 \sigma_{\max} + 0.30 \rho_{\text{unit}} + 0.20 \rho_{\text{subj}} + 0.10 a + 0.05 b, \quad (6)$$

where $\sigma_{\max}$ is the maximum embedding similarity in the retrieved set, ρ_{unit} is the overlap between query key units (entities, dates, numerics) and retrieved tokens, ρ_{subj} is subject-phrase overlap, a is the Jaccard agreement between the conservative and expansive sets, and $b = \min(1, |\text{sources}|/6)$ is a source-diversity bonus.

Routing Decision. The route is selected by comparing the gate scores of the two candidate sets against type-specific thresholds. For *temporal* queries, the conservative set is used by default (*temporal guard*) unless the expansive gate score exceeds $\theta_{\text{temp}} = 0.62$ with agreement ≥ 0.25 , in which case the expansive set is used. For *list*, *number*, *who*, and *where* queries, the *union* of both sets is used when subject overlap is sufficient and gate scores meet minimum thresholds. For *fact* queries, the expansive set is used only when it shows a clear gate-score advantage ($\Delta g \geq 0.08$) over the conservative set; otherwise the conservative set is used. Across 1,986 questions on LoCoMo, this yields 78% conservative, 12% union, and 10% expansive routing.

Rule-Based Answer Verifier. After answer generation, a lightweight verifier checks whether the answer is supported by the retrieved context. The verifier applies type-specific rules without any LLM: for yes/no questions it checks substring support; for list questions it verifies each item against the retrieved evidence; for temporal questions it attempts date-unit recovery from the context. In practice the verifier keeps 97% of answers unchanged, recovers 2%, and abstains on 1%, confirming that its role is conservative post-hoc filtering rather than aggressive rewriting.

3.4 THEORETICAL ANALYSIS

We establish two formal properties that together justify DFM-Fusion’s design and explain its empirical advantage on fact-centric benchmarks. The first (Lemma 3.1) shows that coverage maximization is a submodular optimization problem, justifying the greedy sentence selection in Stage 3. The second (Theorem 3.3) identifies the structural condition under which verbatim packing is provably equivalent to LLM-guided rewriting in terms of answer recall.

Coverage maximization as submodular optimization. Let $S = \bigcup_{i=1}^k \text{sent}(m_i)$ be the full sentence pool after Stage 1, and let $T = T_{\text{num}} \cup T_{\text{ent}} \cup T_{\text{tfidf}} \cup T_{\text{ng}}$ be the salient token set assembled in Stage 4. Define the coverage function

$$f(A) = \frac{|T \cap \bigcup_{s \in A} \text{tok}(s)|}{|T|}, \quad A \subseteq S, \quad (7)$$

where $\text{tok}(s)$ denotes the token multiset of sentence s . The DFM-Fusion selection problem is

$$\max_{A \subseteq S} f(A) \quad \text{s.t.} \quad \sum_{s \in A} |s|_{\text{tok}} \leq B_{\text{fuse}}. \quad (8)$$

Lemma 3.1 (Submodularity of Coverage). *The function f defined in equation 7 is monotone submodular.*

Proof. Monotonicity: For $A \subseteq A'$, $\bigcup_{s \in A} \text{tok}(s) \subseteq \bigcup_{s \in A'} \text{tok}(s)$, so $f(A) \leq f(A')$. *Diminishing returns:* For any $A \subseteq B \subseteq S$ and $s \notin B$, the marginal gain satisfies

$$f(B \cup \{s\}) - f(B) = \frac{|T \cap \text{tok}(s) \setminus \bigcup_{t \in B} \text{tok}(t)|}{|T|} \leq \frac{|T \cap \text{tok}(s) \setminus \bigcup_{t \in A} \text{tok}(t)|}{|T|} = f(A \cup \{s\}) - f(A),$$

since $A \subseteq B$ implies $\bigcup_{t \in A} \text{tok}(t) \subseteq \bigcup_{t \in B} \text{tok}(t)$, so the tokens excluded from s 's marginal contribution are a superset under B . $\square$

Because f is monotone submodular and the budget constraint is a modular knapsack, a greedy algorithm that iteratively selects the sentence with the highest marginal coverage-per-token achieves a $(1 - 1/e)$ -approximation to the optimal solution of equation 8 (Nemhauser et al., 1978; Sviridenko, 2004). Stage 3 maximizes a strength-weighted generalization of f via greedy MMR; the submodular structure is preserved because the strength weights enter as non-negative scalars on the marginal coverage terms.

When verbatim packing is sufficient.

Definition 3.2 (Verbatim Sufficiency). A QA instance (q, a^*) is *Verbatim Sufficient* (VS) with respect to cluster $\mathcal{M}$ if every token of the gold answer a^* appears verbatim in at least one sentence of the pool:

$$\forall w \in \text{tok}(a^*), \quad \exists s \in \mathcal{S} : w \in \text{tok}(s).$$

Theorem 3.3 (Recall Dominance under VS). *Let (q, a^*) be VS with respect to $\mathcal{M}$, and let $s_f = \text{concat}(A)$ be the DFM-Fusion output such that $\text{tok}(a^*) \subseteq \text{tok}(s_f)$ (the answer-bearing sentences are included within budget). Then for any rewriting-based fusion output s_r of $\mathcal{M}$:*

$$\text{Recall}(s_f, a^*) = 1 \geq \text{Recall}(s_r, a^*),$$

where $\text{Recall}(s, a^*) = |\text{tok}(s) \cap \text{tok}(a^*)| / |\text{tok}(a^*)|$.

Proof. By assumption, $\text{tok}(a^*) \subseteq \text{tok}(s_f)$, so $|\text{tok}(s_f) \cap \text{tok}(a^*)| = |\text{tok}(a^*)|$, giving $\text{Recall}(s_f, a^*) = 1$. Since $\text{Recall} \in [0, 1]$, the inequality $1 \geq \text{Recall}(s_r, a^*)$ is immediate for any s_r . $\square$

Theorem 3.3 establishes that under VS, DFM-Fusion achieves perfect answer-token recall—the maximum possible—making LLM-guided rewriting redundant for recall on fact-centric questions. Any LLM rewrite that paraphrases or omits answer tokens strictly reduces recall below 1; any rewrite that preserves them achieves recall equal to DFM-Fusion. The condition “answer-bearing sentences within budget” is satisfied when the budget B_{fuse} is large enough to accommodate the salient high-strength members that carry the answer, as our Stage 3 MMR selection prioritizes.

Remark 3.4. VS holds for fact-centric questions whose answers are directly stated in the conversation: entity names, precise dates, and numerical facts. It fails when the gold answer requires synthesis not present verbatim in the cluster — for example, normalizing “she joined three months after moving” into a relative temporal expression. This boundary is consistent with DFM-Fusion’s observed advantages on multi-hop and single-hop categories and LLM-Fusion’s retained advantages on temporal and open-domain categories (Table 1).

Remark 3.5 (Coverage Certificate vs. LLM Preservation Check). Stage 4 provides a *certified* lower bound: $f(A) \geq \theta_{\text{cov}}$ guarantees that at least θ_{cov} fraction of salient tokens appear verbatim in s_f . By contrast, LLM-Fusion’s preservation check produces an uncertified judgment that may hallucinate confirmation or miss subtle omissions. The certificate is weaker in scope (token-level coverage, not semantic equivalence) but strictly stronger in reliability: it is deterministic and auditable.

3.5 HYPERPARAMETER SELECTION

FadeMem decay parameters ($\theta_{\text{fusion}} = 0.75$, $T_{\text{window}} = 10$, $\lambda_{\text{base}} = 0.1$) are adopted directly from the original system. DFM-specific fusion parameters were initialized to interpretable defaults ($\theta_{\text{dup}} = 0.90$, $B_{\text{fuse}} = 512$, $\lambda = 0.5$, $K = 10$, $\theta_{\text{cov}} = 0.80$) and revised once, in a single diagnostic pass that identified four concrete issues: an embedding representation bug (weighted-average rather than fused-text encoding), a suboptimal text separator, an over-aggressive deduplication threshold, and an undersized fusion budget. This pass was conducted on the full LoCoMo-10 evaluation set without a held-out split. While the changes are individually motivated by interpretable diagnostics rather than metric search, we acknowledge that we cannot fully rule out mild overfitting to the 10 evaluation conversations; the six-model cross-family replication (Section 4) provides indirect stability evidence.

4 EXPERIMENTS

We evaluate DFM-Fusion on LoCoMo to test whether deterministic fusion can replace LLM-guided fusion as the default memory-maintenance operator under a fixed retrieval and answer-generation pipeline.

4.1 EXPERIMENTAL SETUP

Dataset. We use LoCoMo-10 (Maharana et al., 2024), a subset of the LoCoMo benchmark containing 10 conversations with 1,986 QA pairs. Questions span five categories: multi-hop (282), single-hop (817), temporal (321), open-domain (96), and adversarial (446). Multi-hop F1 serves as our primary metric, as it requires integrating information across multiple memory entries and is most sensitive to fusion quality.

Cross-model protocol. The main empirical evaluation is a six-model family study. We run the same memory pipeline across six open-weight answer-generation models served locally via Ollama: qwen25_7b, qwen25_14b, qwen3_8b, llama31_8b, llama32_3b, and mistral_7b. The answer-generation LLM is the sole experimental variable; the memory pipeline (Sentence-BERT embeddings (Reimers & Gurevych, 2019), identical retrieval budgets, and DFM-Fusion or LLM-Fusion operator) is held constant across models. LLM-Fusion maintenance calls (fusion generation and preservation checking) always use DeepSeek-V3.2 regardless of which local answer model runs, replicating FadeMem’s API-based maintenance design. For each model we compare three conditions: (1) **LLM-Fusion**: FadeMem-style fusion with additional DeepSeek-V3.2 calls for fusion generation and preservation checking; (2) **DFM-Fusion**: our deterministic fusion operator with no maintenance-time LLM calls; and (3) **No-Fusion**: no memory merging. All conditions share identical retrieval budgets, scoring protocol, and memory infrastructure.

An additional single-model pilot using DeepSeek-V3.2 as the answer model (3 runs per condition, temperature=0) validated the harness and is reported in Appendix A (Table 9) as a sanity check only; it is not part of the main results.

Statistical tests. We report mean $\pm$ std over 3 runs per condition. Our primary paired comparison is DFM-Fusion vs. LLM-Fusion on per-question multi-hop F1 scores. We report mean differences, paired t-test p -values, and TOST equivalence outcomes with a pre-specified equivalence margin of ± 0.01 absolute multi-hop F1. This margin corresponds to a practically negligible change relative to the observed No-Fusion to LLM-Fusion gaps on LoCoMo-10. The resulting evidence should be interpreted conservatively: only one model satisfies the configured TOST equivalence criterion, so our main claim is practical competitiveness rather than universal statistical equivalence.

4.2 MAIN RESULTS

Accuracy results. Table 1 shows that DFM-Fusion improves over LLM-Fusion on multi-hop F1 in 5/6 models; the only exception is Qwen2.5-7B, where the gap is negligible (-0.0007). The strongest individual result is Qwen3-8B, where DFM-Fusion reaches 0.2157, outperforming LLM-Fusion by $+0.0100$. Averaged across the six families, DFM-Fusion leads by $+0.0036$ on multi-hop F1. The key driver of this advantage is verbatim evidence packing within the retrieval budget, not LLM rewriting — a mechanism that operates identically across model families, which is why the directional result holds consistently despite the heterogeneity of the six answer models.

We observed that the per-category breakdown makes the scope of this advantage precise. DFM-Fusion leads on multi-hop and single-hop categories, where success rewards redundancy-aware packing of verbatim evidence. LLM-Fusion retains small advantages on temporal, open-domain, and adversarial categories, where success requires paraphrastic synthesis or coarse temporal normalization — capabilities that verbatim packing cannot replicate. To illustrate where this boundary lies: on a temporal question such as “When did Caroline join a new activist group?”, DFM-Fusion retrieved topically relevant evidence but could not normalize it into the expected relative temporal expression, returning no date; LLM-Fusion produced a coarse but partially correct relative date. This failure mode is consistent with the average-row temporal gap (DFM 0.0545 vs. LLM 0.0571) and marks exactly the category where generative rewriting retains value.

Table 1: Comprehensive results on LoCoMo-10. Per-model rows show DFM-Fusion (DFM), LLM-Fusion (LLM), and No-Fusion (No) across all six question categories; Δ is DFM minus LLM on multi-hop F1; p is the paired t-test p -value. The **Average** row reports six-model means; bolded cells indicate the best-performing method per column. Ablation rows (*indented*) show DFM variant averages; *italicized* LLM and No values repeat the fixed baselines from the Average row (LLM-Fusion and No-Fusion are unaffected by DFM ablations); $\Delta^\dagger$ compares the ablation to DFM-Fusion on multi-hop F1. DFM-Fusion eliminates all maintenance-time LLM calls: $69.13\times$ average maintenance speedup, 11.74% total-call reduction, 3.88% cost savings. Cost figures use DeepSeek-V3.2 pricing (\$0.27/1M input, \$1.10/1M output); local Ollama models incur zero direct API cost.

Model / Method	Multi-hop F1			Overall F1			Temporal F1			Open-domain F1			Single-hop F1			Adversarial F1			Δ	p
	DFM	LLM	No	DFM	LLM	No	DFM	LLM	No	DFM	LLM	No	DFM	LLM	No	DFM	LLM	No		
Qwen2.5-7B	0.1832	0.1839	0.1769	0.3703	0.3709	0.3644	0.0536	0.0588	0.0536	0.0792	0.0785	0.0739	0.3269	0.3238	0.3133	0.8610	0.8655	0.8655	-0.0007	0.9142
Qwen2.5-14B	0.1836	0.1833	0.1697	0.3714	0.3728	0.3682	0.0505	0.0531	0.0538	0.1216	0.1217	0.1170	0.2830	0.2841	0.2817	0.9417	0.9439	0.9372	+0.0003	0.9510
Qwen3-8B	0.2157	0.2057	0.1992	0.3683	0.3663	0.3643	0.0575	0.0619	0.0610	0.1229	0.1274	0.1394	0.3334	0.3317	0.3262	0.8072	0.8034	0.8072	+0.0100	0.1493
Llama3.1-8B	0.1909	0.1881	0.1825	0.3674	0.3650	0.3611	0.0566	0.0560	0.0580	0.1030	0.1092	0.1108	0.3137	0.3079	0.3085	0.8610	0.8617	0.8453	+0.0028	0.5896
Llama3.2-3B	0.1795	0.1737	0.1780	0.3348	0.3321	0.3337	0.0536	0.0534	0.0494	0.0941	0.0983	0.0838	0.2676	0.2640	0.2694	0.8139	0.8117	0.8117	+0.0058	0.0698
Mistral-7B	0.1731	0.1699	0.1644	0.3383	0.3395	0.3354	0.0549	0.0596	0.0599	0.1046	0.1047	0.1015	0.2907	0.2867	0.2852	0.7870	0.7982	0.7870	+0.0032	0.5526
Average	0.1877	0.1841	0.1785	0.3584	0.3577	0.3545	0.0545	0.0571	0.0560	0.1042	0.1066	0.1044	0.3026	0.2997	0.2974	0.8453	0.8474	0.8423	+0.0036	—
DFM w/o budget truncation	0.1875	<i>0.1841</i>	<i>0.1785</i>	0.3584	0.3577	0.3545	0.0545	0.0571	0.0560	0.1042	0.1066	0.1044	0.3025	0.2997	0.2974	0.8453	<i>0.8474</i>	<i>0.8423</i>	-0.0002 [†]	—
DFM w/ basic coverage	0.1871	<i>0.1841</i>	0.1785	0.3580	0.3577	0.3545	0.0545	0.0571	0.0560	0.1037	0.1066	0.1044	0.3018	0.2997	0.2974	0.8453	<i>0.8474</i>	<i>0.8423</i>	-0.0006 [†]	—
DFM w/o coverage gate	0.1867	<i>0.1841</i>	0.1785	0.3581	0.3577	0.3545	0.0557	0.0571	0.0560	0.1042	0.1066	0.1044	0.3008	0.2997	0.2974	0.8468	<i>0.8474</i>	<i>0.8423</i>	-0.0010 [†]	—
DFM w/ mean embedding	0.1826	<i>0.1841</i>	0.1785	0.3577	0.3577	0.3545	0.0579	0.0571	0.0560	0.1035	0.1066	0.1044	0.3009	0.2997	0.2974	0.8460	<i>0.8474</i>	<i>0.8423</i>	-0.0051 [†]	—

Among the ablation variants, DFM w/ mean embedding is the weakest on multi-hop F1 (-0.0051 vs. DFM-Fusion) while showing the highest temporal F1 (0.0579). We observed that this reversal is consistent with averaged embeddings over-smoothing local structure: mean embeddings retrieve broader, more topically diffuse evidence that occasionally matches temporal queries but fails on the precise multi-hop integration that exact-text embeddings support. The remaining three ablations stay within 0.0010 of DFM-Fusion on multi-hop F1 and differ negligibly across all categories.

Retrieval and fusion diagnostics. Retrieval-side diagnostics support the same direction. Averaged across six models, DFM-Fusion yields slightly stronger evidence recall / hit rate / MRR (0.3667 / 0.4105 / 0.2067) than LLM-Fusion (0.3662 / 0.4098 / 0.2064) and No-Fusion (0.3628 / 0.4070 / 0.2045). Each DFM fusion event packs on average 9.65 selected sentences into 106.64 tokens, with mean coverage of 0.9189 for capitalized entities, 0.8278 for TF-IDF salient terms, 0.8862 for stemmed n -grams, and 1.0 for both dates and numeric tokens. Retrieval diagnostics and fusion behavior statistics (Tables 4–5) are identical or near-identical across the six answer models by design: these quantities depend only on the memory pipeline and the shared Sentence-BERT encoder (Reimers & Gurevych, 2019), not on which local model generates answers.

Retrieval-time redundancy baseline. To isolate whether the observed gains stem from maintenance-time consolidation or could equally be obtained by query-time diversity control, we evaluated **No-Fusion+QueryMMR**: query-time MMR reranking applied at every inference call, without any memory fusion. Averaged across six models in the same experimental run, No-Fusion+QueryMMR achieves 0.1934 multi-hop F1 — exceeding both DFM-Fusion (0.1885) and LLM-Fusion (0.1828) within the same run — confirming that deduplication at any stage of the pipeline yields substantial gains. However, No-Fusion+QueryMMR degrades adversarial F1 by -0.029 relative to DFM-Fusion (0.8160 vs. 0.8453), because per-query diversity reranking introduces noise on adversarial questions that reward selective non-answers. DFM-Fusion achieves comparable multi-hop gains through offline maintenance (a one-time cost amortized over all future queries) while preserving adversarial F1.

Operational efficiency. DFM-Fusion eliminates all maintenance-time LLM calls, replacing them with CPU-only deterministic operations. This yields a $69.13\times$ average maintenance speedup, an 11.74% reduction in total API calls, and a 3.88% reduction in cost — at no accuracy cost on the primary metric. Per-model operational breakdowns, retrieval diagnostics, paired tests, TOST outcomes, and targeted error analysis appear in the appendix.

Table 2: DFM-Route evaluation on LoCoMo-10 across six local models and DeepSeek-V3.2 Pro. DFM-Route (R), DFM-Extractive v2 (E), and No-Fusion+AnchorMMR (N) are compared within the same pipeline. Bold indicates best per column within each model row. DFM-Route uses zero LLM calls at maintenance or retrieval time. **Average** row: six local-model means; DeepSeek-V3.2 row is a single-model API result.

Model	Multi-hop F1			Overall F1			Temporal F1			Single-hop F1			Adversarial F1		
	R	E	N	R	E	N	R	E	N	R	E	N	R	E	N
Qwen2.5-7B	0.1775	0.1468	0.1334	0.4280	0.3662	0.3544	0.3052	0.1936	0.1896	0.3898	0.2702	0.2709	0.8251	0.8767	0.8386
Qwen2.5-14B	0.1745	0.1312	0.1447	0.4591	0.3771	0.3846	0.2962	0.1779	0.1672	0.4023	0.2610	0.2842	0.9462	0.9641	0.9529
Qwen3-8B	0.1977	0.1698	0.1706	0.4423	0.3782	0.3732	0.2950	0.1952	0.1958	0.4467	0.3155	0.3250	0.7683	0.8221	0.7825
Llama3.1-8B	0.2105	0.1767	0.1883	0.4354	0.3719	0.3661	0.2778	0.1663	0.1512	0.4331	0.3161	0.3298	0.7668	0.8117	0.7623
Llama3.2-3B	0.1786	0.1395	0.1591	0.3879	0.3323	0.3345	0.3010	0.1872	0.1843	0.3787	0.2777	0.3009	0.6667	0.7197	0.6771
Mistral-7B	0.1665	0.1417	0.1637	0.4225	0.3568	0.3554	0.3128	0.2126	0.2012	0.4261	0.3025	0.3174	0.7309	0.7623	0.7265
Average	0.1842	0.1509	0.1600	0.4292	0.3637	0.3614	0.2980	0.1888	0.1816	0.4128	0.2905	0.3047	0.7840	0.8261	0.7900
DeepSeek-V3.2 Pro	0.2023	0.1600	0.1510	0.4969	0.4114	0.4101	0.3257	0.2038	0.1985	0.4776	0.3309	0.3364	0.9357	0.9469	0.9410

4.3 DFM-ROUTE EVALUATION

We evaluate DFM-Route (Section 3.3) on the same six model families plus DeepSeek-V3.2 Pro (served via SiliconFlow API), using the `dfm.fusion.optimized.extractive.v2` pipeline as the within-study baseline and `no_fusion.anchor.then_mmr.extractive.v2` as the no-fusion reference. All three conditions share identical memory maintenance (DFM-Fusion), retrieval budgets, and answer-generation stacks; DFM-Route differs only in the retrieval-routing layer at inference time. **Important:** this evaluation uses a different pipeline configuration from the six-model study in Table 1; numbers are not directly comparable across tables.

Accuracy results. Table 2 shows that DFM-Route leads on multi-hop F1 in 6/6 local models and on temporal F1 in 6/6 models. The most striking gain is on temporal F1: the six-model average rises from 0.1888 (DFM-Extractive) to 0.2980 (DFM-Route), a +58% relative improvement. This directly addresses the temporal limitation identified in Section 4: by routing temporal queries to conservative anchor-based retrieval and applying date-unit recovery in the verifier, DFM-Route recovers the temporal precision that verbatim packing alone could not achieve.

The overall F1 improvement ($0.3637 \rightarrow 0.4292$, +0.0655) is driven jointly by temporal and single-hop gains. With DeepSeek-V3.2 Pro as the answer model, DFM-Route reaches 0.4969 overall F1 — the highest result across all experiments.

A consistent trade-off appears on adversarial F1: DFM-Route averages 0.7840 versus 0.8261 for DFM-Extractive. However, we observed that this adversarial gap is identical in character to the adversarial gap of No-Fusion+QueryMMR (Section 4) relative to DFM-Fusion: expanding the retrieval set introduces more potential for incorrectly supported answers, which adversarial questions penalize. The adversarial difference is concentrated in models with weak abstention behavior (Llama3.2-3B, Mistral-7B); Qwen2.5-14B retains near-parity (0.9462 vs. 0.9641).

Routing statistics. Averaged across 1,986 queries on LoCoMo, DFM-Route routes 78% of queries to the conservative strategy, 12% to union, and 10% to expansive. Temporal queries are the main beneficiary of the temporal guard: among the 292 temporal questions, the anchor-based conservative strategy is selected by default, ensuring that date-bearing sentences from the top memory anchor are always included. The rule-based verifier keeps 97% of answers unchanged (1,928/1,986), recovers 2% from over-abstention, and abstains on 1%, confirming that it acts as a conservative post-hoc filter rather than an aggressive rewriter.

5 DISCUSSION AND LIMITATIONS

No universal equivalence claim. The revised evidence does not support a blanket statement that DFM-Fusion is statistically equivalent to LLM-Fusion everywhere. Under the configured TOST test, only Qwen2.5-14B satisfies equivalence. The correct claim is therefore practical rather than ab-

solute: DFM-Fusion is competitive or better on the primary multi-hop metric across a heterogeneous model family, with consistent operational savings.

Benchmark scope. All results are on LoCoMo, which is an appropriate benchmark for fact-centric long-horizon conversational QA but not a sufficient basis for claiming that deterministic fusion replaces LLM-guided consolidation in every setting. The six-model cross-family replication provides evidence of robustness across answer-model families, but not across benchmarks.

Serving-stack dependence. Quality results are model-level, but latency and maintenance numbers are partly serving-stack dependent. The exact speedup factors reported here should therefore be interpreted as deployment-relevant evidence rather than universal constants. What appears robust is the directional result: eliminating maintenance-time fusion calls substantially simplifies and accelerates the memory-maintenance loop.

Surface-biased preservation heuristics. Our deterministic preservation check is based on surface proxies rather than deep semantic entailment. The enriched configuration captures capitalized entities, dates and numerics, TF-IDF salient tokens, and stemmed n -gram coverage, yet it remains English- and formatting-dependent. This leaves open failure modes involving implicit paraphrases, noisy conversational logs, or multilingual settings.

Fallback asymmetry with LLM-Fusion. DFM-Fusion’s coverage gate, when triggered, falls back to concatenating the highest-strength originals within budget — a fused memory is always produced. LLM-Fusion has no equivalent fallback: a failed preservation check discards the entire cluster, leaving memories unmerged. In practice this asymmetry is moot for the reported experiments (LLM-Fusion rejected 0/113 clusters across all six models), but it means DFM-Fusion is more conservative at the preservation gate while still guaranteeing a consolidation result.

Embedding model sensitivity. All reported experiments use Sentence-BERT (Reimers & Gurevych, 2019) for both sentence deduplication and retrieval. Cross-encoder sensitivity was not part of the main result. We provide runnable reviewer-study configurations for E5 and BGE-style sentence encoders, but treat those as sensitivity checks rather than a new headline claim.

Single-benchmark scope and benchmark generalization. All results are on LoCoMo, which is appropriate for fact-centric long-horizon conversational QA. Tasks requiring abstraction, preference synthesis, or explicit reconciliation of conflicting evidence may benefit more from generative maintenance operators, as the targeted error analysis already shows (temporal normalization, open-domain preference synthesis). Evaluation on additional benchmarks such as LongMemEval is an important next step.

Query-time redundancy control. We implemented No-Fusion+QueryMMR as a retrieval-time-only baseline: query-time MMR without maintenance fusion. Averaged across six models, this achieves 0.1934 multi-hop F1, exceeding DFM-Fusion (0.1885) within the same run. This result confirms that the deduplication mechanism — not LLM rewriting — is the primary driver of performance gains, whether applied at maintenance or retrieval time. The trade-off is a -0.029 drop in adversarial F1 (0.8160 vs. 0.8453 for DFM-Fusion), because per-query diversity reranking introduces noise on questions that reward selective non-answers. DFM-Fusion achieves comparable multi-hop gains through offline maintenance (a one-time cost amortized over all future queries), while DFM-Route extends this at retrieval time with query-type-aware routing, recovering the temporal precision that query-agnostic deduplication cannot address (temporal F1 +58%; Section 4.3).

Why gap recovery is no longer the right headline. In the earlier single-setting draft, gap recovery was a convenient scalar because the ordering between No-Fusion and LLM-Fusion was clean. In the six-model family this is less stable; for example, Llama3.2-3B places No-Fusion slightly above LLM-Fusion on multi-hop F1. For this reason, we emphasize direct DFM–LLM deltas, model-count wins (5/6), and cross-method averages rather than using gap recovery as the primary summary statistic.

6 CONCLUSION

The central finding of this paper is that the gain attributed to LLM-guided memory fusion in fact-centric long-horizon conversational QA comes almost entirely from deduplicating and packing verbatim evidence into a fixed retrieval budget — not from the LLM’s linguistic rewriting ability. DFM-Fusion, a fully deterministic four-stage operator requiring zero maintenance-time LLM calls, recovers this gain across a heterogeneous family of six open-weight models. DFM-Route, its query-aware retrieval extension, further demonstrates that the remaining temporal gap is a retrieval-routing problem rather than a rewriting problem: by routing temporal queries to anchor-based retrieval with date-unit recovery, temporal F1 improves $5.5\times$ (0.055 to 0.298) without any LLM call at maintenance or retrieval time. Together, DFM-Fusion and DFM-Route demonstrate that the full memory pipeline — from consolidation to evidence delivery — can operate without any API dependency in the maintenance loop.

At the same time, our results do not support a blanket equivalence claim. DFM-Fusion’s core advantage is specific to fact-centric QA, where the evaluation rewards verbatim evidence retrieval; open-domain preference synthesis may still benefit from paraphrastic rewriting. An adversarial trade-off accompanies the expanded retrieval in DFM-Route: routing more queries to union or expansive strategies introduces noise that adversarial questions penalize, a tension that warrants further study. Future work should extend this investigation to benchmarks requiring abstraction and explicit conflict reconciliation, and explore hybrid operators that selectively invoke LLM rewriting only when coverage verification signals a genuine synthesis gap. Limitations of the current study — including single-benchmark scope, surface-biased preservation heuristics, and the absence of a held-out validation split — are discussed in Section 5.

A ADDITIONAL RESULTS

A.1 PER-MODEL OPERATIONAL BREAKDOWN

Table 3: Per-model latency, call volume, and cost breakdown. Cost figures use DeepSeek-V3.2 pricing (\$0.27/1M input, \$1.10/1M output) applied to all LLM calls. Local Ollama answer models incur zero direct API cost; their cost figures reflect maintenance-call overhead only. DFM-Fusion cost is identical across all six models (zero fusion calls, same QA call count). Answer-generation cost estimated at 1,585 input + 30 output tokens per QA call.

Model	LLM calls	DFM calls	LLM fusion calls	LLM maint.	DFM maint.	LLM total	DFM total	LLM cost	DFM cost
Qwen2.5-7B	2212	1986	226	332.6	5.5	458.2	272.2	0.948	0.916
Qwen2.5-14B	2292	1986	306	635.0	5.7	1018.4	705.3	0.960	0.916
Qwen3-8B	2212	1986	226	423.0	9.5	618.2	384.2	0.948	0.916
Llama3.1-8B	2264	1986	278	473.3	4.3	893.8	637.2	0.955	0.916
Llama3.2-3B	2280	1986	294	393.4	7.7	618.2	456.9	0.957	0.916
Mistral-7B	2244	1986	258	296.8	7.6	695.7	542.0	0.952	0.916

A.2 RETRIEVAL DIAGNOSTICS

Table 4: Per-model retrieval quality diagnostics. Retrieval statistics depend solely on the memory pipeline and the shared Sentence-BERT encoder (Reimers & Gurevych, 2019), not on the answer-generation model; identical or near-identical values across the six answer models are expected by design.

Model	DFM Rec.	LLM Rec.	No Rec.	DFM Hit	LLM Hit	No Hit	DFM MRR	LLM MRR	No MRR
Qwen2.5-7B	0.3667	0.3657	0.3628	0.4105	0.4093	0.4070	0.2067	0.2063	0.2045
Qwen2.5-14B	0.3667	0.3667	0.3628	0.4105	0.4103	0.4070	0.2067	0.2067	0.2045
Qwen3-8B	0.3667	0.3657	0.3628	0.4105	0.4093	0.4070	0.2067	0.2063	0.2045
Llama3.1-8B	0.3667	0.3676	0.3628	0.4105	0.4112	0.4070	0.2067	0.2066	0.2045
Llama3.2-3B	0.3667	0.3657	0.3628	0.4105	0.4093	0.4070	0.2067	0.2064	0.2045
Mistral-7B	0.3667	0.3657	0.3628	0.4105	0.4093	0.4070	0.2067	0.2063	0.2045

Table 5: Per-model DFM-Fusion behavior and coverage diagnostics. DFM-Fusion is fully deterministic; fusion behavior statistics are independent of the answer-generation model and identical across models sharing the same fusion-event history.

Model	Cluster	Selected	Tokens	Fallback	Cap.	TF-IDF	Stem	Date	Num.
Qwen2.5-7B	3.4592	9.6735	106.8878	46	0.9189	0.8301	0.8852	1.0000	1.0000
Qwen2.5-14B	3.4592	9.6735	106.8878	46	0.9189	0.8301	0.8852	1.0000	1.0000
Qwen3-8B	3.4490	9.6224	106.3878	44	0.9189	0.8255	0.8872	1.0000	1.0000
Llama3.1-8B	3.4592	9.6735	106.8878	46	0.9189	0.8301	0.8852	1.0000	1.0000
Llama3.2-3B	3.4490	9.6224	106.3878	44	0.9189	0.8255	0.8872	1.0000	1.0000
Mistral-7B	3.4490	9.6224	106.3878	44	0.9189	0.8255	0.8872	1.0000	1.0000

A.3 FUSION BEHAVIOR

A.4 TOST OUTCOMES

Table 6: Per-model TOST equivalence outcomes for DFM-Fusion vs. LLM-Fusion on multi-hop F1.

Model	DFM-LLM delta	<i>p</i> -value	TOST equivalent
Qwen2.5-7B	-0.0007	0.9142	false
Qwen2.5-14B	+0.0003	0.9510	true
Qwen3-8B	+0.0100	0.1493	false
Llama3.1-8B	+0.0028	0.5896	false
Llama3.2-3B	+0.0058	0.0698	false
Mistral-7B	+0.0032	0.5526	false

A.5 TARGETED ERROR ANALYSIS

Across the inspected model logs, DFM-Fusion underperforms LLM-Fusion in two recurring ways. The dominant failure mode is temporal retrieval miss: the retrieved evidence is either absent or incomplete, after which LLM-Fusion sometimes still answers with a coarse relative date such as “last Tuesday” or “last week.” A second, smaller failure mode appears in open-domain questions that require paraphrastic synthesis rather than verbatim evidence packing. These patterns match the category-level averages in Table 1 (lower block), where LLM-Fusion retains small advantages on temporal and open-domain questions.

Table 7: Representative targeted error cases extracted from model-level analyses.

Model	Failure mode	Representative question	Interpretation
Qwen2.5-7B	answer generation mismatch	Which pet did Jolene adopt first, Susie or Seraphim?	Retrieved evidence was partial, and DFM answered “Not mentioned” while LLM-Fusion returned the correct short answer.
Qwen3-8B	temporal retrieval miss	When did Caroline join a new activist group?	DFM returned no date after a retrieval miss, while LLM-Fusion produced a coarse but partially correct relative date.
Qwen3-8B	temporal normalization	When did Evan meet his future wife?	DFM preserved topical evidence (Canada trip) but failed to normalize it into the expected relative temporal expression.
Llama3.1-8B	paraphrase synthesis or	Which type of vacation would Evan prefer, city walking tours or camping outdoors?	The open-domain answer requires synthesizing preference evidence into a short decision, a setting where LLM-Fusion still has an advantage.

A.6 DFM-ROUTE PER-MODEL RESULTS

Table 8: DFM-Route (R), DFM-Extractive v2 (E), and No-Fusion+AnchorMMR (N) per-model results on LoCoMo-10 (3 runs each, mean reported). This table uses the `extractive_v2` pipeline; numbers are not directly comparable to Table 1 which uses a different pipeline configuration. DeepSeek-V3.2 Pro is a single-model API result (SiliconFlow).

Model	Multi-hop F1			Overall F1			Temporal F1			Single-hop F1			Adversarial F1		
	R	E	N	R	E	N	R	E	N	R	E	N	R	E	N
Qwen2.5-7B	0.1775	0.1468	0.1334	0.4280	0.3662	0.3544	0.3052	0.1936	0.1896	0.3898	0.2702	0.2709	0.8251	0.8767	0.8386
Qwen2.5-14B	0.1745	0.1312	0.1447	0.4591	0.3771	0.3846	0.2962	0.1779	0.1672	0.4023	0.2610	0.2842	0.9462	0.9641	0.9529
Qwen3-8B	0.1977	0.1698	0.1706	0.4423	0.3782	0.3732	0.2950	0.1952	0.1958	0.4467	0.3155	0.3250	0.7683	0.8221	0.7825
Llama3.1-8B	0.2105	0.1767	0.1883	0.4354	0.3719	0.3661	0.2778	0.1663	0.1512	0.4331	0.3161	0.3298	0.7668	0.8117	0.7623
Llama3.2-3B	0.1786	0.1395	0.1591	0.3879	0.3323	0.3345	0.3010	0.1872	0.1843	0.3787	0.2777	0.3009	0.6667	0.7197	0.6771
Mistral-7B	0.1665	0.1417	0.1637	0.4225	0.3568	0.3554	0.3128	0.2126	0.2012	0.4261	0.3025	0.3174	0.7309	0.7623	0.7265
Average	0.1842	0.1509	0.1600	0.4292	0.3637	0.3614	0.2980	0.1888	0.1816	0.4128	0.2905	0.3047	0.7840	0.8261	0.7900
DeepSeek-V3.2 Pro	0.2023	0.1600	0.1510	0.4969	0.4114	0.4101	0.3257	0.2038	0.1985	0.4776	0.3309	0.3364	0.9357	0.9469	0.9410

A.7 DEEPSEEK-V3.2 SINGLE-MODEL SANITY CHECK

Before running the six-model family study, we validated the harness using DeepSeek-V3.2 (API, temperature=0) as the answer model, with 3 independent runs per condition. These results confirmed the directional ordering (DFM-Fusion > No-Fusion, DFM-Fusion $\approx$ LLM-Fusion) and are included only as a harness sanity check; the main empirical basis of the paper is the six-model local-model evaluation in Section 4.

Table 9: DeepSeek-V3.2 single-model sanity check (3 runs per condition, temperature=0). Reported as harness validation only.

Setting	Multi-hop F1	Overall F1
LLM-Fusion	0.1863 $\pm$ 0.0021	0.3805 $\pm$ 0.0004
DFM-Fusion	0.1872 $\pm$ 0.0034	0.3795 $\pm$ 0.0012
No-Fusion	0.1710 $\pm$ 0.0009	0.3750 $\pm$ 0.0011

REFERENCES

- Jaime Carbonell and Jade Goldstein. The use of mmr, diversity-based reranking for reordering documents and producing summaries. In *Proceedings of the 21st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 335–336, 1998.
- Prateek Chhikara, Dev Khant, Saket Aryan, Taranjeet Singh, and Deshraj Yadav. Mem0: Building production-ready ai agents with scalable long-term memory, 2025.
- Darren Edge, Ha Trinh, Newman Cheng, Joshua Bradley, Alex Chao, Apurva N. Mody, Steven Truitt, and Jonathan Larson. From local to global: A graph rag approach to query-focused summarization. *ArXiv*, abs/2404.16130, 2024.
- Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Qianyu Guo, Meng Wang, and Haofen Wang. Retrieval-augmented generation for large language models: A survey. *ArXiv*, abs/2312.10997, 2023.
- Bernal Jimenez Gutierrez, Yiheng Shu, Yu Gu, Michihiro Yasunaga, and Yu Su. Hipporag: Neurobiologically inspired long-term memory for large language models. *ArXiv*, abs/2405.14831, 2024.
- Chuanrui Hu, Xingze Gao, Zuyi Zhou, Dannong Xu, Yi Bai, Xintong Li, Hui Zhang, Tong Li, Chong Zhang, Lidong Bing, and Yafeng Deng. Evermemos: A self-organizing memory operating system for structured long-horizon reasoning. *ArXiv*, abs/2601.02163, 2026.

-
- Kai Li, Xuanqing Yu, Ziyi Ni, Yi Zeng, Yao Xu, Zheqing Zhang, Xin Li, Jitao Sang, Xiaogang Duan, Xuelei Wang, Chengbao Liu, and Jie Tan. Timem: Temporal-hierarchical memory consolidation for long-horizon conversational agents. *ArXiv*, abs/2601.02845, 2026.
- Jiaqi Liu, Yaofeng Su, Peng Xia, Siwei Han, Zeyu Zheng, Cihang Xie, Mingyu Ding, and Huaxiu Yao. Simplemem: Efficient lifelong memory for llm agents. 2026.
- Adyasha Maharana, Dong-Ho Lee, S. Tulyakov, Mohit Bansal, Francesco Barbieri, and Yuwei Fang. Evaluating very long-term conversational memory of llm agents. *ArXiv*, abs/2402.17753, 2024.
- George L. Nemhauser, Laurence A. Wolsey, and Marshall L. Fisher. An analysis of approximations for maximizing submodular set functions—I. *Mathematical Programming*, 14(1):265–294, 1978.
- Charles Packer, Vivian Fang, Shishir G. Patil, Kevin Lin, Sarah Wooders, and Joseph Gonzalez. Memgpt: Towards llms as operating systems. *ArXiv*, abs/2310.08560, 2023.
- Nils Reimers and Iryna Gurevych. Sentence-bert: Sentence embeddings using siamese bert-networks. *ArXiv*, abs/1908.10084, 2019.
- Maxim Sviridenko. A note on maximizing a submodular set function subject to a knapsack constraint. *Operations Research Letters*, 32(1):41–43, 2004.
- Mohammad Tavakoli, Alireza Salemi, Carrie Ye, Mohamed Abdalla, Hamed Zamani, and J. R. Mitchell. Beyond a million tokens: Benchmarking and enhancing long-term memory in llms. *ArXiv*, abs/2510.27246, 2025.
- Lei Wei, Xiao Peng, Xu Dong, Niantao Xie, and Bin Wang. Fademem: Biologically-inspired forgetting for efficient agent memory. 2026.